

Programming Paradigms

Lecture 11

Slides are from Prof. Chin Wei-Ngan from NUS

More on Declarative Concurrency

Agents and Message Passing Concurrency


Ports

- A **port** is an ADT with two operations:
 - $\{\text{NewPort } S \ P\}$ or equivalently $P = \{\text{NewPort } S\}$: create a new port with entry point (channel) P and stream S .
 - $\{\text{Send } P \ x\}$: append x to the stream corresponding to the entry point P .
- Successive sends from the same thread appear on the stream in the same order in which they were executed.
- This property implies that a port is an asynchronous FIFO (first-in, first-out) communication channel.

Semantics of Ports

- Extend the execution state of the declarative model by adding a **mutable store** μ
- This store contains ports, i.e. pairs of the form $x : y$, where x and y are variables of the single-assignment store (x is the channel's name and y is the *current last position of stream*).
- The mutable store is initially empty.
- The semantics guarantees that x is *always bound* to a name value that represents a port and that y is *unbound*.
- The execution state becomes a triple (MST, σ, μ) (or (MST, σ, μ, τ) if the trigger store is considered).

The Message-Passing Concurrent Mode



The NewPort Operation

- The semantics of $(\{\text{NewPort } \langle x \rangle \langle y \rangle\}, E)$ is:
 - Create a **fresh port name** (also called **unique address**) n .
 - Bind $E(\langle y \rangle)$ and n in the store.
 - If the binding is successful, then add the pair $E(\langle y \rangle) : E(\langle x \rangle)$ to the mutable store μ .
 - If the binding fails, then raise an error condition.

The Send Operation

- The semantics of $(\{\text{Send } \langle x \rangle \langle y \rangle\}, E)$ is:
 - If the activation condition is true ($E(\langle x \rangle)$ is determined), then:
 - If $E(\langle x \rangle)$ is not bound to the name of a port, then raise an error condition.
 - If the mutable store contains $E(\langle x \rangle) : z$, then:
 - Create a new variable $z0$ in the store.
 - Update the mutable store to be $E(\langle x \rangle) : z0$.
 - Create a new list pair $E(\langle y \rangle) | z0$ and bind z with it in the store.
 - If the activation condition is false, then suspend execution.

Question

```
declare S P  
P={NewPort S}  
{Browse S}  
thread {Send P a} end  
thread {Send P b} end
```

- What will the Browser show?
- Note that each `{Send P ...}` is in a separate thread

Question

```
declare S P  
P={NewPort S}  
{Browse S}  
thread {Send P a} end  
thread {Send P b} end
```

- Which will the Browser show?
- Either
 - $a \mid b \mid _ <future>$ or
 - $b \mid a \mid _ <future>$
- non-determinism: we can't say what

Answering Messages

- Traditional view
- Include the entry port P' of the sender in the message:
 $\{\text{Send } P \text{ pair(Message } P')\}$
- Receiver sends answer message to P'
 $\{\text{Send } P' \text{ AnsMessage}\}$

Answering Messages

- Do not reply by address, use something like pre-addressed reply envelope
 - dataflow variable!!!
- {Send P pair(Message Answer) }
- Receiver can bind Answer!

Port Objects

- A **port object** is a combination of one or more ports and a stream object.
- This *extends stream objects* in two ways:
 - First, many-to-one communication is possible: many threads can reference a given port object and send to it independently.
 - This is not possible with a stream object because it has to know where its next message will come from.
 - Second, port objects can be embedded inside data structures (including messages).
 - This is not possible with a stream object because it is referenced by a stream that can be extended by just one thread.

Stream Object

```
input           accumulator           output
      ↓             ↓                   ↑
proc {StreamObject S1 X1 ?T1}
    case S1 of M|S2 then N X2 T2 in
        {NextState M X1 N X2}
        T1 = N|T2 {StreamObject S2 X2 T2}
    [] nil then T1=nil end
end

declare S0 X0 T0
thread {StreamObject S0 X0 T0} end
```

```
StreamObject :: [A], B, [C] → ()
NextState :: A,B, C,A → ()
```

Port Objects. Distributed Algorithm

```
declare P1 P2 ... Pn in
local S1 S2 ... Sn in
  {NewPort S1 P1}
  {NewPort S2 P2}
  ...
  {NewPort Sn Pn}
thread {RP S1 S2 ... Sn} end
end
```

- The thread contains a recursive procedure RP that reads the port streams and performs some action for each message received.
- Sending a message to the port object is just sending a message to one of the ports.
- Similar terms: **agent**, **process** (Erlang), **active object**

A Math Agent

```
proc {Math E}  
    case E  
        of add(N M Answer) then Answer=N+M  
        [] mul(N M Answer) then Answer=N*M  
        [] int(Formula Answer) then  
            Answer = ...  
    end  
end
```

- **Remark:** Answer is included in the stream's element
 x of {Send EntryPoint X}

Making the Agent Work (Port Creation)

```
local S in
    MP = {NewPort S}
    proc {MathProcess Ms}
        case Ms of M|Mr then
            {Math M}
            {MathProcess Mr}
        end
    end
    thread {MathProcess S} end
end
```

- `MathProcess` is a recursive procedure that reads the port streams and performs some action for each message received.

Making the Agent Work (Sending a Message)

```
declare A B  
thread % client 1  
  {Send MP add(2 3 A)}  
  {Browse A}  
end  
thread % client 2  
  {Send MP mul(2 3 B)}  
  {Browse B}  
end
```

- A and B are two dataflow variables which will be bound in port MP

Recall Higher-Order Construct

```
ForAll :: {[X], X→()} → ()
```

```
proc {ForAll Xs P}  
  case Xs  
    of nil  then skip  
    [] X|Xr then {P X} {ForAll Xr P}  
  end  
end
```

- Call procedure P for all elements in Xs

Smells of Higher-Order...

- Using `ForAll`, we have

```
proc {MathProcess Ms}
  {ForAll Ms Math}
end
```

Making the Agent Work

```
declare MP in
local S in
  MP = {NewPort S}
  thread {ForAll S Math} end
end
```

Making the Agent Work

```
declare MP in
local S in
  MP = {NewPort S}
  thread for M in S do {Math M} end end
end
```

- The stream `s` is private (`local`) to the port.
- `Math` is associated to the port `MP`
- `MP` and `Math` can become arguments of a generic function.

Smells Even Stronger...

- Programming with port objects can be abstracted into a function

```
fun {NewAgent Process}
```

```
    Port Stream
```

```
NewAgent :: {X→()} → Port X
```

```
in
```

```
    Port={NewPort Stream}
```

```
    thread {ForAll Stream Process} end
```

```
    Port
```

```
end
```

- So, the previous port creation is equivalent with:

```
MP = {NewAgent Math}
```

Why Do Agents/ Processes Matter?

- Model to capture communicating entities
- Each agent is simply defined in terms of how it replies to messages
- Each agent has a thread of its own
 - no screw-up with concurrency
 - we can easily extend the model so that each agent has a state (encapsulated)
- ***Extremely useful to model distributed systems!***

Summary so far

- Ports for message sending
 - use stream (list of messages) as mailbox
 - port serves as unique address
- Use agent abstraction
 - combines port with thread running agent
 - simple concurrency scheme
- Introduces non-determinism... and state!

Protocols

Protocols

- **Protocol:** is a set of rules for sending and receiving messages
 - programming with agents
- Most well-known protocols:
 - the Internet protocols (TCP/IP, HTTP, FTP, etc.)
 - LAN (Local Area Network) protocols such as Ethernet and DHCP (Dynamic Host Connection Protocol), ...

RMI (Remote Method Invocation)

- It seems to be the most popular of the simple protocols.
- It allows an object to call another object in a different operating system process, either on the same machine or on another machine connected by a network.
- **RMI** is a descendant of the **RPC** (Remote Procedure Call), which was invented in 1980, before object-oriented programming became popular.
- RMI became popular once objects started replacing procedures as the remote entities to be called.
- We assume that a “*method*” is simply what a port object does when it receives a particular message.

Differences between RPC and RMI


RPC

- Faster than RMI
- Depends on the platform
- Has to convert the arguments between architectures so that each computer can use its native datatype

RMI

- Is part of Java's object-oriented approach
- Allows multiple-concurrent method invocation
- Is portable (doesn't depend on the operating system)
- Good security system
- To call outside methods, RMI needs JNI, JDBC, RMI-IIOP, RMI-IDL, etc.

1. RMI (2 calls)



- A client sends a request to a server and then waits for the server to send back a reply.
- C stands for client, S for server, *idle* means “available to service requests”, *suspended* means “not available”.

The Server as a Port Object

```
declare
proc {ServerProc Msg}
    case Msg
    of calc(X Y) then
        Y = X * X + 1.0
    end
end
Server={NewAgent ServerProc}
```

- The second argument `Y` of `calc` is bound by the server.
- The server computes the polynomial `x * x + 1.0`

What is NewAgent? (Reminder)

```
fun {NewAgent Process}
    Port Stream
in
    Port={NewPort Stream}
    thread {ForAll Stream Process} end
    Port
end
```

The Client (using RMI)

```
declare
proc {ClientProc Msg}
  case Msg
    of work(Y) then
      Y1 Y2 in
        {Send Server calc(1.0 Y1)}
        {Wait Y1}
        {Send Server calc(2.0 Y2)}
        {Wait Y2}
      Y = Y1 + Y2
    end
  end
Client={NewAgent ClientProc}
```

The Client as a Port Object II

```
local X in
    {Send Client work(X)}
    {Browse X}
end
```

- Difference between the client and server:
 - The client definition references the server directly but the server definition does not know its clients.
 - The server gets a client reference indirectly, through the argument `Y`, i.e. the dataflow variable that is bound to the answer by the server.
 - The client waits until receiving the reply before continuing.

What is Wait?

- `{Wait X}` suspends the thread until `X` becomes determined, i.e. also called *explicit synchronization* on variable `X`

declare `Y`

`{ByNeed proc { $ X } X=1 end Y}`

`{Browse Y}`

`{Wait Y}`

`<statement>`

- Display `Y` in the Browser.
- To access `Y`, the operation `{Wait Y}` will trigger the producing procedure.
- `<statement>` will be executed only after `Y` is bound


Characteristics of RMI

- In the previous example, all messages are executed sequentially by the server.
- In practice, some RMI implementations do things somewhat differently, i.e. they allow multiple calls from different clients to be processed concurrently.
- May use different languages and different OS.

Asynchronous RMI

- Similar to RMI, except that the client continues execution immediately after sending the request.
- The client is informed when the reply arrives.
- So, two requests can be done in rapid succession.
- **Motivation:** If communications between client and server are slow, then this will give a large performance advantage over RMI.

2. Asynchronous RMI *(2 calls)*



The Asynchronous RMI Client

```
declare
proc {ClientProc Msg}
  case Msg
  of work(Y) then Y1 Y2 in
    {Send Server calc(1.0 Y1)}
    {Send Server calc(2.0 Y2)}
    Y = Y1 + Y2
  end
end
Client={NewAgent ClientProc}
local X in
  {Send Client work(X)}
  {Browse X}
end
```

Characteristics of Asynchronous RMI

- Message sends overlap. Client waits for both results y_1 and y_2 before doing the addition y_1+y_2 .
- The server is the same as with standard RMI. It still receives messages one by one and executes them sequentially.
- Requests are handled by the server in the same order as they are sent and the replies arrive in that order as well.

RMI with Callback

- Server may need to call back client to fulfill request, e.g. check on some special values.

```
proc {ServerProc Msg}
    case Msg
        of calc(X ?Y Client) then X1 D in
            {Send Client delta(D)} ← callback
            X1=X+D
            Y = X * X + 1.0
        end
    end
Server={NewAgent ServerProc}
```

RMI with Callback

```
proc {ClientProc Msg}
    case Msg
        of work(?Z) then Y in
            {Send Server calc(10.0 Y Client)}
            Z=Y+100.0
        [] delta(?D) then D=1.0
    end
end
Client={NewAgent ClientProc}
{Browse {Send Client work($)} }
```

- Does this work?


No! It deadlocks as server and client waiting for each other.

Solution – Use Thread

```
proc {ClientProc Msg}  
    case Msg  
        of work(?Z) then Y in  
            {Send Server calc(10.0 Y Client)  
             thread Z=Y+100.0 end  
             [] delta(?D) then D=1.0  
        end  
    end
```

add thread to allow client to proceed.

RMI with Callback (using thread)



RMI with Callback (using continuation)

- Possible to avoid thread.


```
proc {ServerProc Msg}
  case Msg
    of calc(X Client Cont) then X1 D Y in
      {Send Client delta(D)}
      X1=X+D
      Y = X * X + 1.0
      {Send Client Cont#Y}
    end
  end
Server={NewAgent ServerProc}
```

The diagram illustrates the continuation flow in the code. A pink arrow points from the word 'Cont' in the 'calc' message to the continuation block where it is used as a parameter. Another pink arrow points from the continuation block back to the 'Cont' variable, labeled 'continuation'.

Solution – Using Continuation Record

```
proc {ClientProc Msg}
    case Msg
        of work(?Z) then Y in
            {Send Server calc(10.0 Client cont(Z)) }
            [] cont(Z) #Y then Z=Y+100.0
            [] delta(?D) then D=1.0
        end
    end
Client={NewAgent ClientProc}
{Browse {Send Client work($)} }
```

RMI with Callback (using continuation record)



Erlang

Erlang

- Developed by Ericsson for telecoms application.
- Features : fine grain parallelism, extreme reliability, hot code updates.
- Functional core – dynamically typed strict functional language.
- Message-passing extension – processes communicate by sending messages asynchronously in FIFO order.

Functions in Erlang

- Uses pattern-matching and Prolog syntax

```
factorial(0) -> 1;
```

```
factorial(N) when N>0 -> N * factorial(N-1).
```

Pattern-Matching with Tuple

```
area({square, Side}) -> Side*Side;
```

```
area({rectangle,X,Y}) -> X*Y;
```

```
area({circle, R}) -> 3.14159*R*R;
```

```
area({triangle, A,B,C}) -> ... ;
```

tuple

Concurrency and Message Passing

- `spawn(M, F, A)` creates a new process and returns its `Pid`. Note that M-module, F-initial function, A-argument list.
- Send operation (written as `Pid!msg`) is an asynchronous message sending.
- `receive` operation removes message from a mailbox. It uses pattern-matching to select messages for removal

An Erlang Process

```
-module(areaserver)
-export([start/0, loop/0])

start() -> spawn(areaserver, loop, []).

loop() -> receive
            {From, Shape} ->
                From!area(Shape),
                loop()
        end.
```

spawn

receive

send

Receive Construct

```
receive
    Pattern1 [when Guard1] -> Body1;
    :
    PatternN [when GuardN] -> BodyN;
    [after Expr -> BodyT; ]
end
```

This expression blocks until a message matching
one of patterns arrives or when timeout occurs

Summary

- Stream Object
- Thread Module and Composition
- Soft Real-Time Programming
- Agents and Message Passing
- Protocols
- Erlang